

Project Report

Credit Card Default Prediction

Using Classification and Risk-Based Techniques

Submitted to : Finance Club, IIT Roorkee

Name:	Jay Vaghasiya
Branch:	CSE 2nd Year
Enrollment:	23114104
Email:	vaghasiya_jr@cs.iitr.ac.in

IIT Roorkee

Contents

1	Project Overview	3
1.1	What is Credit Card Default and Why is it Important?	3
1.2	Business Objective	3
1.3	Given Data	3
1.4	Problem Statement	4
1.5	Approach	4
2	Data Preprocessing	5
3	Exploratory Data Analysis (EDA)	6
3.1	Demographic Analysis	6
3.2	Financial Analysis	7
4	Feature Engineering From Behavioral Analysis	8
4.1	Credit Utilization Ratios	8
4.2	Repayment and Billing Volatility	8
4.3	Delay-Based Features	8
4.4	Payment Trend Features	9
4.5	Billing Pattern Change	9
5	Featur Selection	10
5.1	Feature Importance Analysis	10
6	Model Training and Evaluation	12
6.1	Logistic Regression	12
6.2	Random Forest	12
6.3	LightGBM	13
6.4	XGBoost	13
7	Model Comparison and Selection	14
7.1	Threshold Optimization	14
7.2	Business Implications of False Positives vs. False Negatives	14
7.3	risk Based Model Selection	15
8	Final Predictions and Conclusion	15

1 Project Overview

1.1 What is Credit Card Default and Why is it Important?

Credit card default occurs when a customer is unable to repay their outstanding credit card balance. This poses significant financial risk to banks and lending institutions. Predicting which customers are likely to default in the upcoming month helps minimize losses, manage credit exposure, and take proactive action. Using historical repayment behavior, credit usage, and demographic data, our goal is to build a machine learning model that flags high-risk customers before they default. This enables smarter credit decisions and strengthens financial stability for lenders.

1.2 Business Objective

The primary objective of this project is to assist financial institutions in minimizing credit risk by identifying customers who are likely to default on their credit card payments in the upcoming month. Default risk is a critical factor in credit-based systems, and proactively managing it allows lenders to reduce financial losses, optimize credit allocation, and maintain portfolio health.

To address this, we apply classification-based machine learning techniques combined with risk-aware strategies such as threshold tuning, recall-focused evaluation F score, and feature engineering aligned with financial risk indicators (e.g., credit utilization, delay streaks, payment-to-bill ratios). Our model aims not only to predict defaulters accurately but to prioritize catching high-risk cases, even at the cost of occasional false alarms. This aligns with the core business goal: minimizing missed defaulters while maintaining operational efficiency.

1.3 Given Data

The dataset provided contains historical credit and repayment records for a large number of credit card customers. Each row represents a unique customer, and the columns describe both demographic attributes and financial behavior over a six-month period.

Key columns include:

- **LIMIT_BAL** – The credit limit assigned to the customer.
- **SEX, EDUCATION, MARRIAGE, AGE** – Basic demographic information.
- **BILL_AMT1 to BILL_AMT6** – Monthly bill statements from the past six months.
- **PAY_AMT1 to PAY_AMT6** – Monthly repayment amounts over the same six months.
- **PAY_0 to PAY_6** – Delay status of payments for each month (e.g., -1 = paid in advance, 0 = on time, 1 = one month delay, etc.).
- **next_month_default** – The target variable indicating whether the customer defaulted on payment in the following month (1 = default, 0 = no default).

The training dataset contains labeled records used to train and validate the model, while the final test (validation) set consists of similar customer records without the target variable. Our task is to accurately predict the default status for this validation set.

1.4 Problem Statement

The objective of this project is to develop a machine learning model that can accurately predict whether a credit card customer will default on their payment in the upcoming month. The prediction must be made based on historical customer behavior, including credit utilization, repayment history, and demographic factors.

This is a binary classification task where the goal is to identify high-risk customers before they default. By accurately forecasting defaults, financial institutions can proactively manage credit risk, reduce potential losses, and make more informed lending decisions. The model must be optimized not just for accuracy, but for recall and F_2 -score, in order to prioritize catching defaulters even at the cost of occasional false positives.

1.5 Approach

The solution to the credit card default prediction problem was designed using a structured pipeline of data analysis, transformation, modeling, and evaluation. The approach followed these major steps:

- **Step 1: Data Exploration** – We began with exploratory data analysis to understand the distribution of features, identify patterns between customer attributes and default behavior, and check for class imbalance or missing values.
- **Step 2: Feature Engineering** – Based on domain understanding, we created new features such as credit utilization, payment-to-bill ratios, repayment delay streaks, and payment trends over time. These features aimed to capture customer behavior more effectively than the raw variables.
- **Step 3: Addressing Class Imbalance** – The dataset had far fewer defaulters compared to non-defaulters. We experimented with techniques like SMOTE and class weights to ensure the model paid sufficient attention to the minority class.
- **Step 4: Model Training and Evaluation** – We trained and compared multiple models including Logistic Regression, Random Forest, LightGBM, and XGBoost. Each model was evaluated using validation metrics such as F_2 -score, recall, and AUC-ROC.
- **Step 5: Threshold Optimization** – Since our objective prioritized recall, we tuned the prediction threshold for each model to maximize the F_2 -score, which weights recall higher than precision.
- **Step 6: Feature Selection and Finalization** – Using feature importance metrics from tree-based models, we removed low-impact features to reduce noise and improve generalization.
- **Step 7: Final Prediction** – The best-performing model (XGBoost with tuned threshold and selected features) was used to generate predictions for the unseen test dataset.

2 Data Preprocessing

Before training the models, several preprocessing steps were performed to clean and prepare the data:

- **Missing Values:** A total of 126 missing values were found in the `AGE` column. These were imputed using the median age of the dataset to preserve distribution symmetry and avoid introducing bias.
- **Inconsistent Categorical Entries:** Several categorical columns contained invalid or inconsistent values:
 - Values of 0 in the `marriage` column were replaced with 3 (`Others`), as 0 is not a valid marital status code.
 - Values 0, 5, 6 in the `education` column were also inconsistent. These were grouped and replaced with 4 (`Others`). Since these values accounted for over 1% of the data (311 entries), they were corrected rather than removed.
- **Incorrect Values in Precomputed Columns:** Two columns present in the original dataset contained invalid or inconsistent negative values:
 - `AVG_Bill_amt`: Although this column was provided in the dataset, some values were negative, which is not valid in the context of billing amounts. These incorrect entries were corrected by recalculating the average manually using `BILL_AMT1` to `BILL_AMT6`.
 - `PAY_TO_BILL_ratio`: This column, also present in the dataset, contained some negative values that are not meaningful in a financial context. These were clipped to zero to ensure the ratio remained within a valid range.

3 Exploratory Data Analysis (EDA)

3.1 Demographic Analysis

Demographic features such as marriage, sex, education, and age were analyzed to understand their distribution and relationship with default behavior.

- **marriage:** The dataset contains three categories: 1 (married), 2 (single), and 3 (others). distribution shows that percentage of defaulters is slightly higher for married customers across all three categories.
- **sex:** Plot says that the percentage of defaulters is slightly higher for Female customers compared to male customers.
- **education:** Here education is categorized into 1 (graduate school), 2 (university), 3 (high school), and 4 (others). The distribution shows that the percentage of defaulters is highest for customers with High School education, followed by those with University and then Graduate School education. The "Others" category has the lowest default rate.
- **age:** Initially plotted box-plot of age indicated some outliers, which we handled by clipping with IQR method. Then KDE plot of age against the target variable `next_month_default` shows that the percentage of defaulters is highest for customers in the age group of 25-35 years, and it gradually decreases for older age groups. This suggests that younger customers are more likely to default compared to older customers.




Figure 1: Distributions of Age, Marriage, Sex, and Education against Default Status

3.2 Financial Analysis

Financial features such as credit limit, bill amounts, payment amounts, and payment delays were analyzed to understand their relationship with default behavior.

- **LIMIT_BAL**: KDE plot of credit limit against the target variable `next_month_default` shows that customers with lower credit limits (below 1,00,000) have a higher percentage of defaults compared to those with higher credit limits. Also data is highly skewed towards lower credit limits, inducting outliers with high values but since these can be so meaningful in financial context representing (premium customers), we decided to keep them for further analysis.
- **BILL_AMT1 to BILL_AMT6**: While analyzing the monthly bill amounts, we observed that for each month, defaulters tend to have lower bill amounts compared to non-defaulters. This suggests that defaulters may be using their credit cards less or are unable to manage higher spending.
- **AVG_Bill_amt**: The average bill amount across the six months shows a similar trend, with defaulters having lower average bills compared to non-defaulters and this feature also highly skewed towards left side leading to more than 2000 outliers, but again we decided to keep them as they represent meaningful financial behavior.
- **PAY_AMT1 to PAY_AMT6**: The repayment amounts for defaulters are generally lower than those for non-defaulters. This indicates that defaulters may be making smaller payments or struggling to meet their repayment obligations.
- **PAY_0 to PAY_6**: The payment delay status shows that defaulters tend to have more instances of delayed payments compared to non-defaulters. This suggests that a history of delayed payments is a strong indicator of potential default.
- **PAY_TO_BILL_ratio**: The ratio of payment to bill amount shows that defaulters tend to have lower ratios compared to non-defaulters. This indicates that defaulters may be struggling to keep up with their bills relative to their payments.




Figure 2: Distributions of Financial Features against Default Status

4 Feature Engineering From Behavioral Analysis

While traditional EDA focuses on understanding the structure and distribution of data, financial and behavioral analysis dives deeper into customer repayment patterns, credit usage behavior, and risk indicators that more directly relate to default likelihood. This type of analysis is crucial in credit risk modeling, as default events are often driven by subtle trends in a customer's financial decisions over time. In this section, we evaluated key behavioral signals such as credit utilization,

bill growth, delay streaks, payment consistency, and repayment-to-bill ratios. These insights were instrumental in designing features that aligned with how lending institutions assess credit risk and customer reliability.

4.1 Credit Utilization Ratios

Credit utilization reflects how much of a customer's available credit limit is being used, which is a key indicator in credit risk scoring. We calculated the following utilization-based features:

- **Util_0:** Ratio of the most recent month's bill (`Bill_amt1`) to the credit limit. Indicates immediate credit pressure.
- **UTIL_max:** Highest utilization across six months. Captures peak credit usage and potential financial stress points.
- **UTIL_avg:** Average utilization across six months. Reflects long-term credit behavior.

Higher utilization values, especially peak and average, were commonly observed in defaulters, indicating risk due to over-leveraging credit.

4.2 Repayment and Billing Volatility

To capture consistency in customer financial behavior, we calculated standard deviations for repayment and billing amounts across six months:

- **pay_std:** Measures how much monthly repayments vary. High values indicate irregular payments and financial instability; low values suggest disciplined repayment habits.
- **bill_std:** Measures variation in monthly bills. High values reflect erratic spending behavior, while low values point to steady and controlled usage.

Defaulters generally exhibited higher volatility in both measures, making these features valuable for risk prediction.

4.3 Delay-Based Features

To capture customer payment discipline, we engineered features that quantify both the frequency and pattern of delayed payments:

- **overdue_cnt:** Total number of months with delayed payment (where $\text{delay} \geq 0$). More delays often signal higher default risk.

- **on_time_ratio:** Proportion of months paid on time or in advance. A high ratio indicates good repayment behavior.
- **num_months_late:** Counts the number of months with significant delay ($\text{PAY}_x \geq 1$). Provides another view of repayment discipline.
- **longest_delay_streak:** Maximum number of consecutive months with payment delay. Longer streaks are strongly associated with future defaults.

4.4 Payment Trend Features

Understanding how payments evolve over time helps identify financial stress or improving discipline:

- **payment_slope:** Linear trend (slope) of monthly payments. A negative slope suggests declining repayment ability, which may indicate rising risk.
- **pay_ewm_sum:** Weighted sum of recent payments, giving more weight to recent months. Captures short-term repayment behavior effectively.

4.5 Billing Pattern Change

- **bill_growth_1:** Measures the change in bill amount from month 2 to month 1. A sharp increase may suggest sudden spending spikes, which could precede financial strain.




Figure 3: Behavioral Features Engineered Based on Behavioral risk Based Analysis

5 Featur Selection

After engineering a comprehensive set of features, we performed feature selection to identify the most impactful variables for our classification task. This step is crucial to reduce noise, improve model interpretability, and enhance generalization.

5.1 Feature Importance Analysis

We have analyzed feature importance using correlation coefficients, mutual information scores, and tree-based feature importance metrics. After evaluating the impact of each feature on the target variable `next_month_default`, we identified the best representative key features that significantly contribute to predicting credit card defaults.




Figure 4: Feature Importance from XGBoost Model

- **marriage** – Marital status has a moderate correlation with default risk.
- **LIMIT_BAL** – Credit limit is a strong predictor of financial behavior.
- **AVG_Bill_amt** – Average bill amount reflects overall credit usage.
- **PAY_TO_BILL_ratio** – Indicates how well customers manage their bills relative to payments.
- **longest_delay_streak** – Captures the pattern of delayed payments, which is critical for risk assessment.
- **num_months_late** – Total months with significant delay directly correlates with default likelihood.
- **Util_0, UTIL_avg, UTIL_max** – Credit utilization ratios are strong indicators of financial stress.

- `overdue_cnt` – Number of overdue months is a direct measure of repayment discipline.
- `on_time_ratio` – Proportion of on-time payments indicates reliability.
- `bill_growth_1` – Sudden changes in billing patterns can signal potential issues.
- `bill_std`, `pay_std` – Variability in billing and payment amounts reflects financial stability.
- `pay_ewm_sum` – Recent payment trends provide insights into current financial behavior.

These features were selected based on their strong correlation with the target variable and their ability to capture different aspects of customer behavior that contribute to default risk. when combined, these features provide a comprehensive view of customer financial health, credit usage patterns, and repayment discipline which we have proved through correlation analysis which shows that these features are highly correlated with the target variable `next_month_default`.

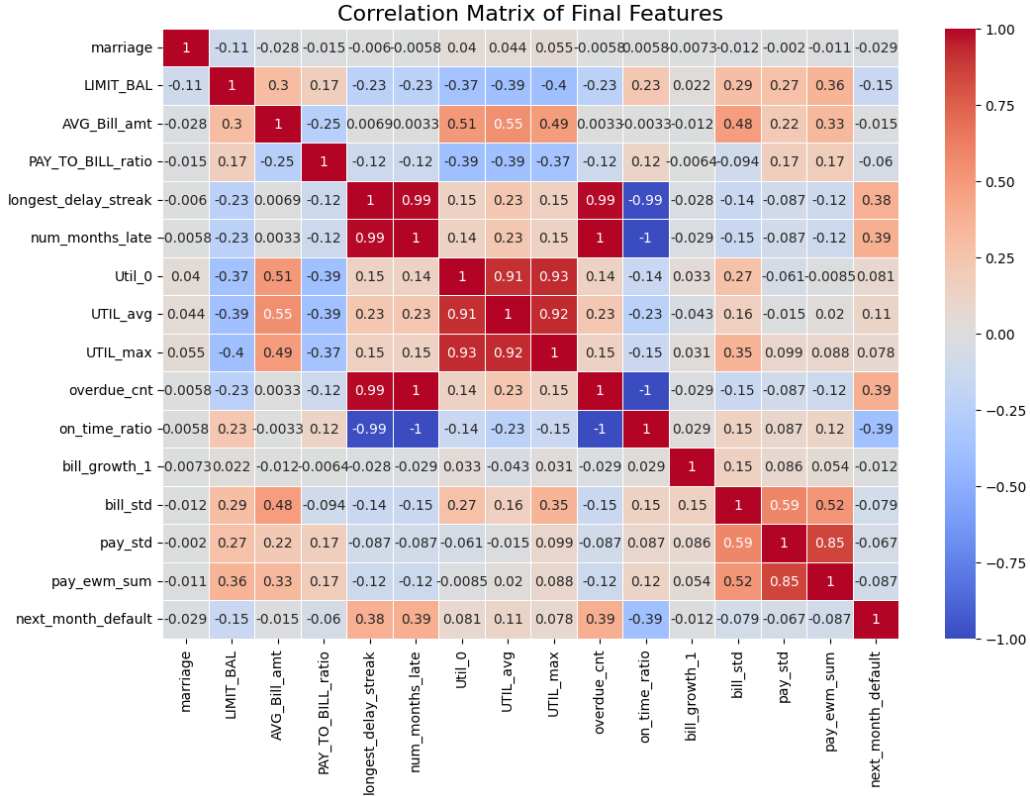


Figure 5: Correlation Heatmap of Selected Features with Target Variable

6 Model Training and Evaluation

Multiple machine learning models were trained and evaluated to predict whether a customer will default in the upcoming month. Each model was tested with different strategies for handling class imbalance and optimizing performance using validation data. This section outlines the models used and compares their results.

6.1 Logistic Regression

Logistic Regression was used as a baseline linear model due to its interpretability and speed. Several variations were tested:

- **Class-Weighted Logistic Regression:** The model was trained from inbuilt standard library `sklearn.linear_model.LogisticRegression` with `class_weight='balanced'` to give more importance to the minority class (defaulters). This approach helped improve recall and F_2 -score compared to standard training.
- **SMOTE + Logistic Regression:** The training set was oversampled using SMOTE to synthetically generate additional defaulter samples. Although recall improved, the overall F_2 -score dropped slightly due to reduced model calibration, making this method less optimal.
- **Logistic Regression with Cross-Validation:** A version using stratified cross-validation was implemented to ensure stable performance across folds. This helped validate the consistency of the model but did not significantly outperform the simpler class-weighted version.

Logistic Regression served as a strong baseline, but its performance plateaued around an F_2 -score of 0.58, leading us to explore non-linear models that could capture more complex relationships in the data.

6.2 Random Forest

Random Forest, a tree-based ensemble model, was tested to capture non-linear patterns and feature interactions that logistic regression could not. Several variations were explored:

- **Class-Weighted Random Forest:** The model was trained with `class_weight='balanced'` to adjust for class imbalance. This improved F_2 -score significantly over logistic regression and provided better recall with strong AUC performance.
- **SMOTE + Random Forest:** Oversampling with SMOTE followed by Random Forest training was tested, but the model's recall dropped sharply and the F_2 -score decreased. This indicated overfitting to synthetic examples, and the approach was discarded.
- **Feature-Optimized Random Forest:** A version using only top-ranked features (based on importance) showed slightly better generalization and a more interpretable model without hurting performance.

Random Forest delivered a strong balance between recall and AUC, achieving an F_2 -score of approximately 0.606, making it a top contender before moving to boosting-based models.

6.3 LightGBM

LightGBM, a gradient boosting framework optimized for speed and efficiency, was trained as a next step. It is particularly well-suited for tabular data and handles categorical-like features and missing values effectively.

- **Baseline LightGBM:** Trained with `class_weight='balanced'` and default hyperparameters. It achieved a strong F_2 -score (0.612) with significantly better recall than Random Forest.
- **SMOTE + LightGBM:** SMOTE oversampling was applied before training. Although accuracy improved, recall dropped, leading to a lower F_2 -score. Hence, the original class-weighted version was preferred.
- **Feature-Pruned LightGBM:** After identifying low-importance features, we removed them and retrained the model. This improved both F_2 -score and recall slightly, while keeping the model simpler and more focused.

LightGBM consistently performed well and remained a strong candidate for final selection due to its speed, recall, and balanced generalization.

6.4 XGBoost

XGBoost, a powerful gradient boosting algorithm known for its regularization and performance on structured data, was selected as the final model due to its consistent and top-tier results across evaluation metrics.

- **Class-Weighted XGBoost:** The model was trained using `scale_pos_weight` to address class imbalance. This configuration yielded high recall and the best F_2 -score among all models (0.6149), with strong AUC performance (0.777).
- **Threshold Optimization:** Default 0.5 threshold was replaced by a custom threshold (0.314), which was selected based on maximum F_2 -score from the precision-recall curve. This tuning significantly improved defaulter detection.
- **Feature Selection:** Low-importance features were removed based on XGBoost's gain-based feature importance. This improved model focus and resulted in slightly better performance.

XGBoost emerged as the best-performing model, achieving an F_2 -score of 0.6149 and a recall of 0.66 on the validation set. Its ability to handle complex interactions and provide robust predictions made it the final choice for credit card default prediction.

7 Model Comparison and Selection

After evaluating multiple models on validation data, we compared their performance using key classification metrics: Accuracy, Recall, F1 Score, F₂-Score, and AUC-ROC. The goal was to first decide which metric to prioritize based on business objectives, and then select the best model accordingly.

Model	Accuracy	Recall	F1 Score	F ₂ -Score	AUC-ROC
Logistic Regression	0.774	0.558	0.485	0.582	0.746
Random Forest	0.784	0.591	0.511	0.606	0.774
LightGBM	0.756	0.652	0.504	0.612	0.775
XGBoost (Final)	0.755	0.649	0.503	0.6149	0.777

Table 1: Model Comparison on split Validation Data

7.1 Threshold Optimization

To align the model’s predictions with our primary evaluation metric (F₂-score), we did not use the default probability cutoff of 0.5. Instead, we plotted the precision–recall curve on the validation set and computed F₂-score at each candidate threshold. The F₂-score is defined as

$$F_2 = \frac{(1 + 2^2) \text{Precision} \times \text{Recall}}{2^2 \text{Precision} + \text{Recall}},$$

which weights recall twice as heavily as precision. We found that a **threshold = 0.314** maximized F₂ for our final XGBoost model, yielding **F₂ = 0.6149**. All final predictions on the unseen validation data were generated using this cutoff.

7.2 Business Implications of False Positives vs. False Negatives

In credit risk management, the cost of a false negative (FN)—failing to flag a true defaulter—is substantially higher than the cost of a false positive (FP)—flagging a non-defaulter as high-risk.

- **False Negatives (Missed Defaulters):** Lead to direct financial losses, increased provisioning requirements, and erosion of risk controls.
- **False Positives (Over-Flags):** Cause inconvenience to customers and minor operational costs (e.g., additional review), but do not incur large credit losses.

By maximizing F₂-score, our model prioritizes recall (minimizing FN) even if it allows more FP. This risk-based threshold selection ensures the bank can proactively intervene with high-risk accounts, preserving portfolio health.

7.3 risk Based Model Selection

In our risk-based evaluation framework, we prioritized recall (minimizing missed defaulters) and F_2 -score above all, leading to the selection of XGBoost as the final model for three key reasons:

1. **Highest F_2 -Score:** At the optimized threshold of 0.314, XGBoost achieved an F_2 of 0.6149, surpassing all other candidates.
2. **Strong Recall and AUC:** With recall = 0.649 and AUC-ROC = 0.777, it balances aggressive defaulter detection with robust ranking capability.
3. **Focused Feature Use:** After pruning low-importance features, XGBoost maintained top performance, indicating strong generalization and alignment with financial risk signals.

These factors confirm XGBoost as the best choice for production-style predictions on the validation dataset, ensuring maximal defaulter coverage with controlled false alarms.

8 Final Predictions and Conclusion

From the final XGBoost model, we generated probability of default for each customer in unseen `validate_dataset_final.csv`. and this is how probability distribution looks like:




Figure 6: Probability Distribution of Default Probabilities

The final predictions were made by applying the optimized threshold of 0.31428874 to the predicted probabilities. Customers with probabilities above this threshold were classified as likely to default (1), while those below were classified as non-defaulters (0). Out of the total 5016 customers in the validation set we found :

Default (1)	No Default (0)
2910	2106

Table 2: Final Predictions on Validation Set

The final model's predictions were saved in file `submission_23114104.csv` in the required format `[Customer_ID, next_month_default]`. In github repository provided below you can find the datasets, model training scripts, and final submission file.

GitHub

Thank You!